The benchmark protocol: 250 questions written before any output, answered by a frozen agent, scored twice

Plan R152-R157, left to right as the data flows. The agent sees only the questions; both scores compare its answers with the gold answers. Blue and orange: the two scores, as in the R157 charts.

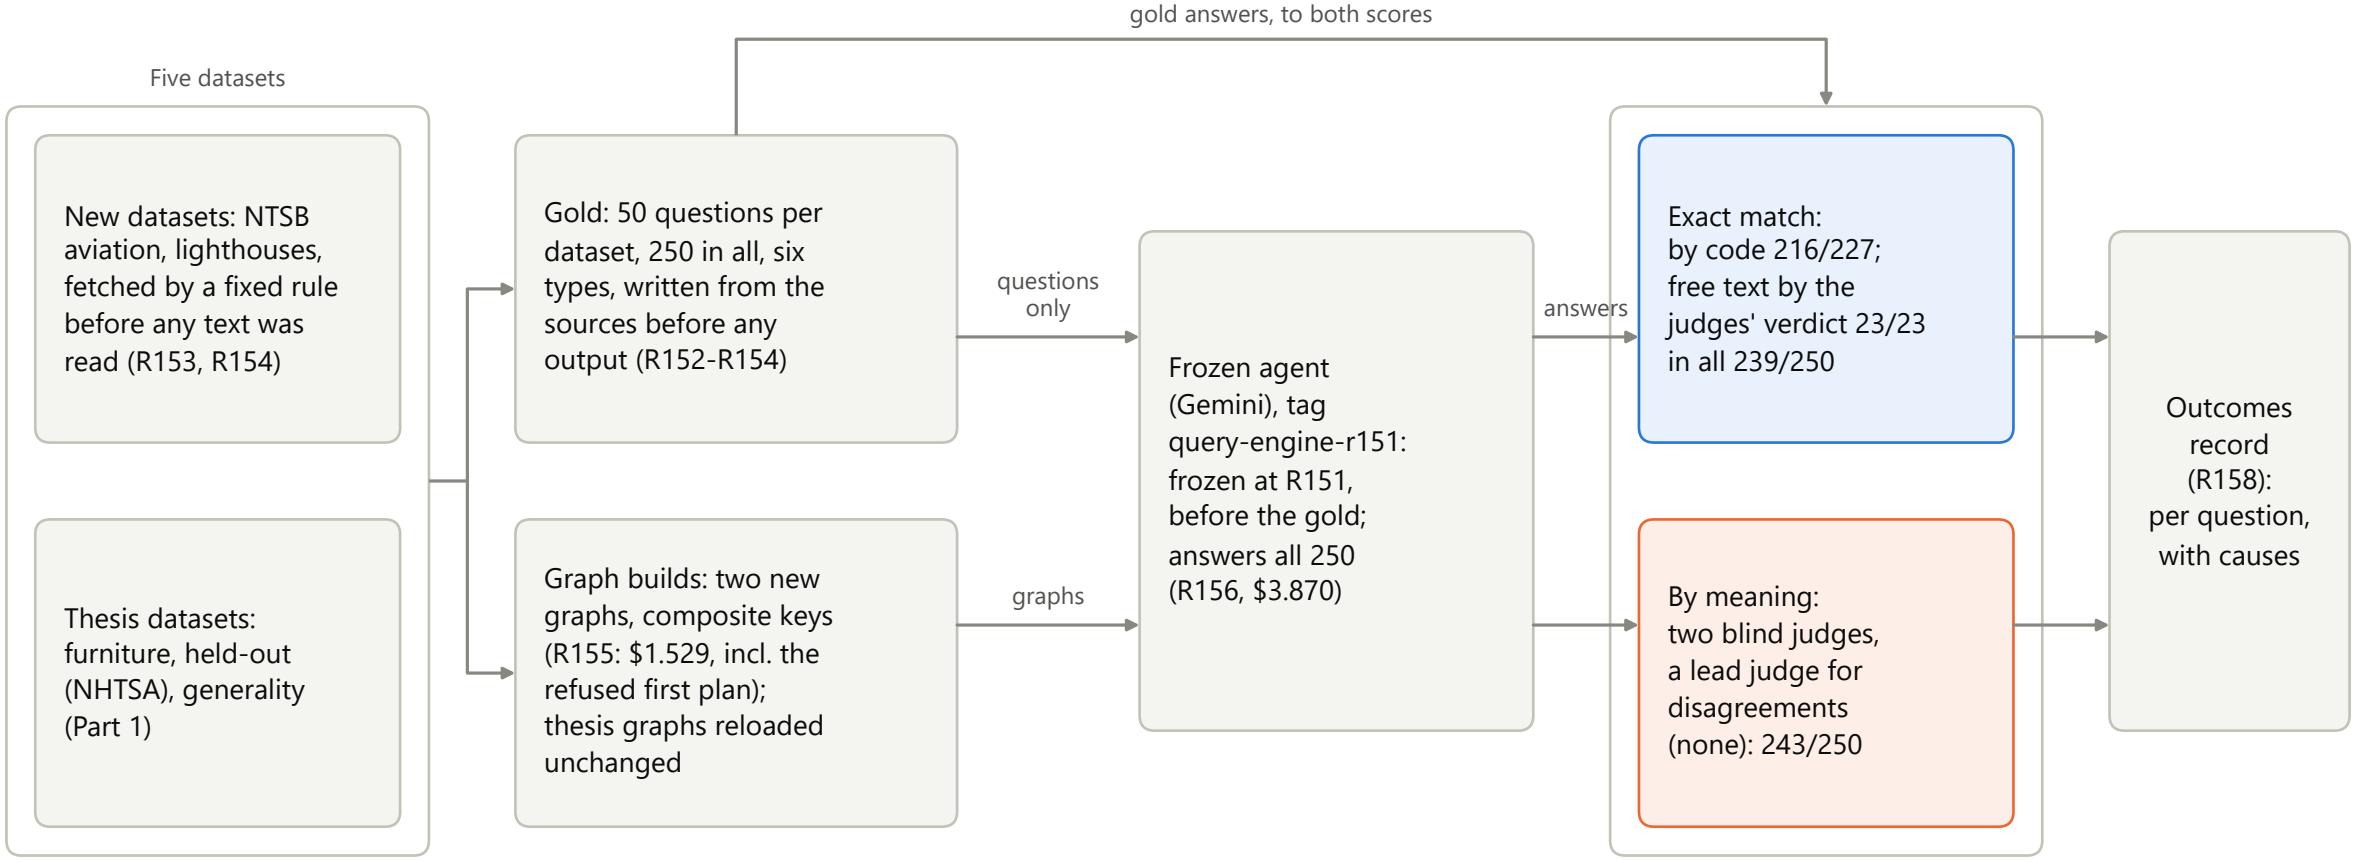

Gold and judges come from one model family (Claude Opus 5.5), which also wrote the agent's fixes; the agent answers with Gemini. One run, no repeat.
Each judge compares the answer with its gold answer, without the other's verdicts.