Measure, fix, measure again: three question sets written blind, the engine fixed between them from its traced failures

Plans R152-R157, R160-R166 and R167-R173, left to right in time. Within a column the questions and the engine feed the run; a column's failures feed only the next column's fixes.

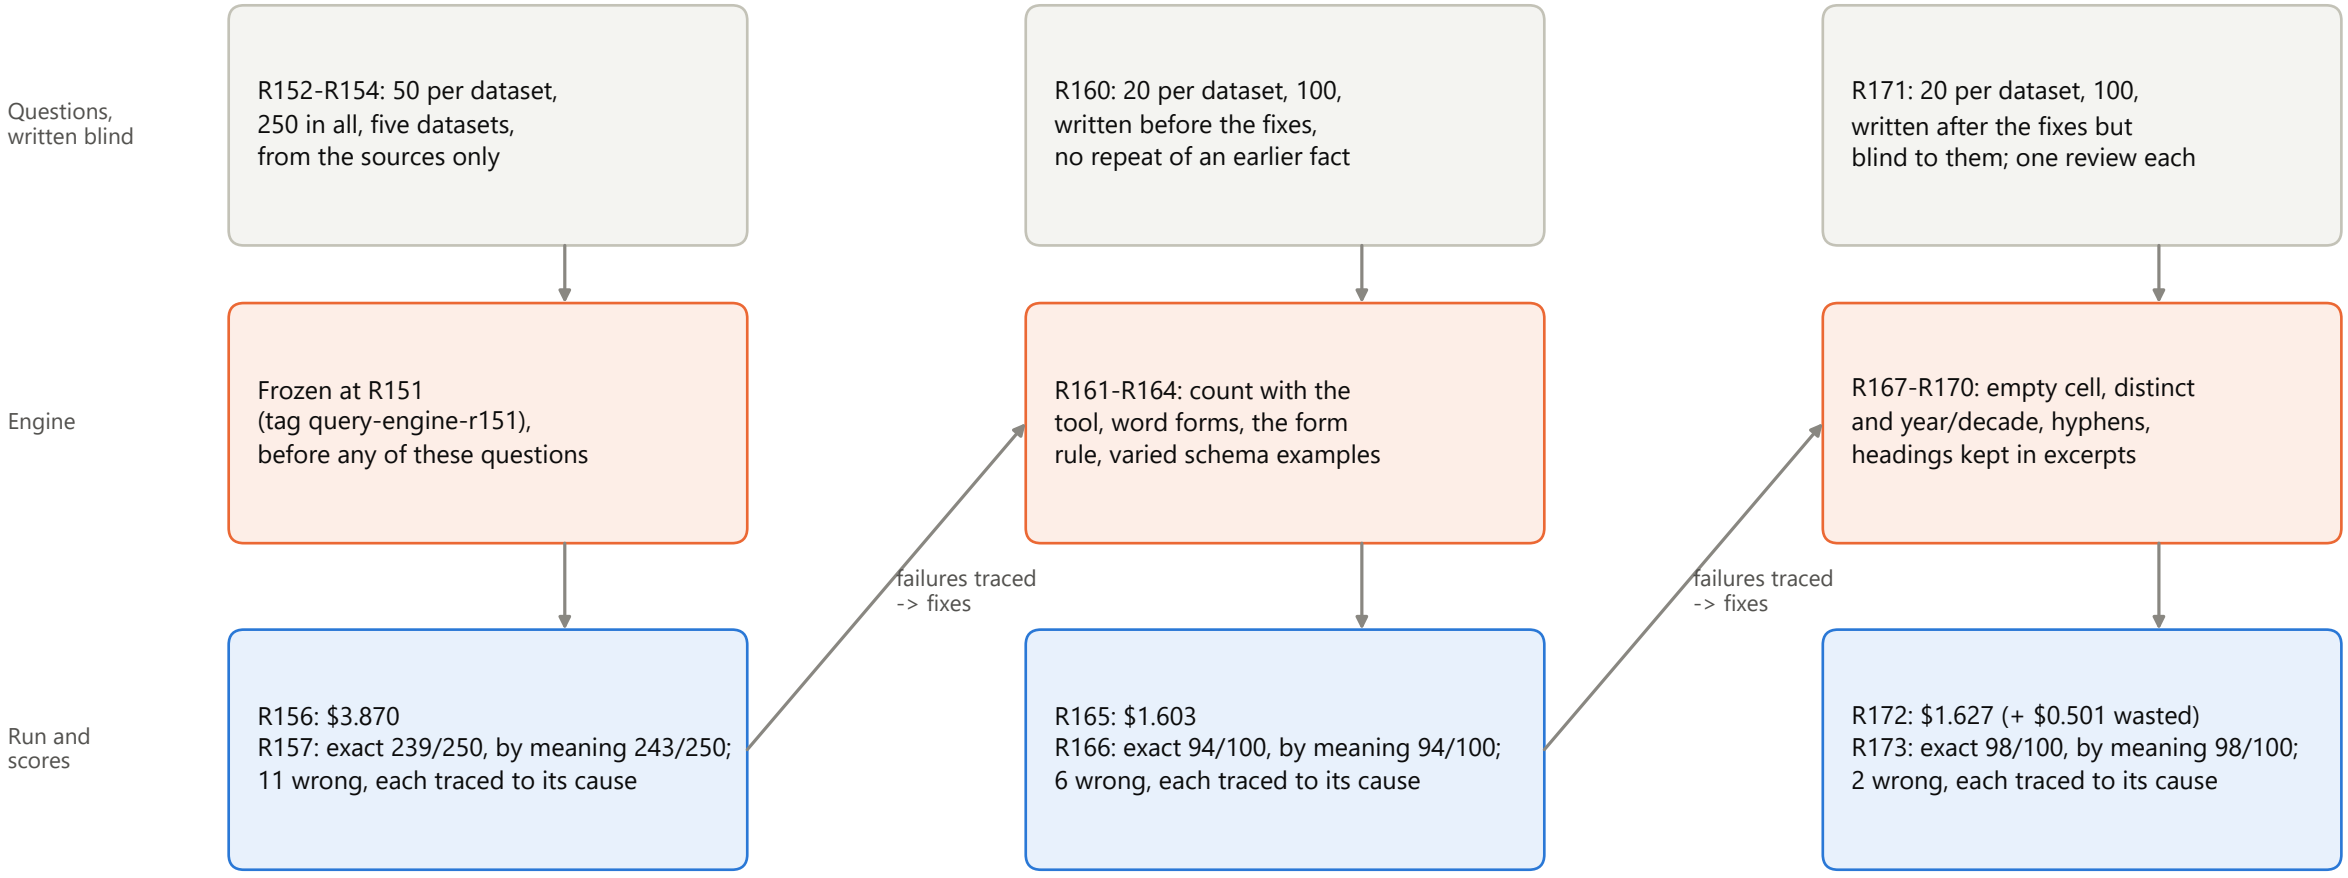

Blue: the scores, as in the charts; orange: the engine. Each fix was proven by an offline replay of the recorded failing calls ($0) before the next paid run.
The tool text changes with every fix, so each score describes the engine of its column on its own questions; the question sets differ, so no difference is attributed to the fixes.