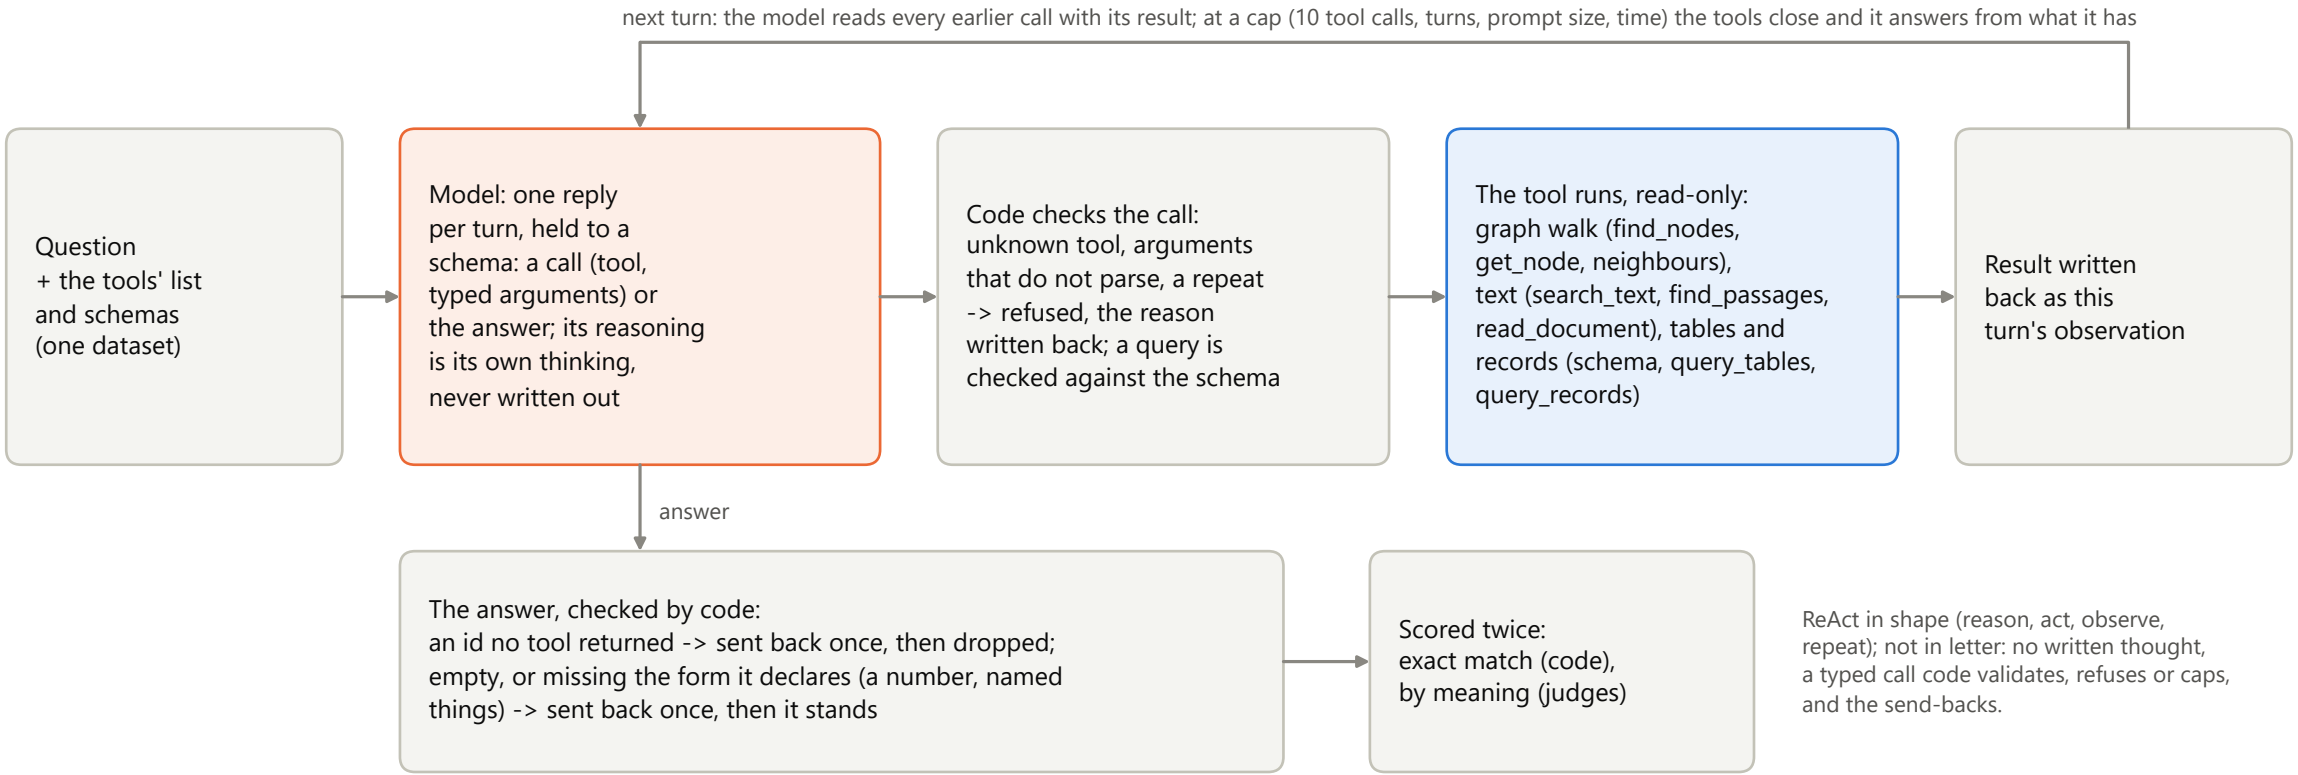
next turn: the model reads every earlier call with its result; at a cap (10 tool calls, turns, prompt size, time) the tools close and it answers from what it has
Question
+ the tools' list
and schemas
(one dataset)
Model: one reply
per turn, held to a
schema: a call (tool,
typed arguments) or
the answer; its reasoning
is its own thinking,
never written out
Code checks the call:
unknown tool, arguments
that do not parse, a repeat
-> refused, the reason
written back; a query is
checked against the schema
The tool runs, read-only:
graph walk (find_nodes,
get_node, neighbours),
text (search_text, find_passages,
read_document), tables and
records (schema, query_tables,
query_records)
Result written
back as this
turn's observation
answer
The answer, checked by code:
an id no tool returned -> sent back once, then dropped;
empty, or missing the form it declares (a number, named
things) -> sent back once, then it stands
Scored twice:
exact match (code),
by meaning (judges)
ReAct in shape (reason, act, observe,
repeat); not in letter: no written thought,
a typed call code validates, refuses or caps,
and the send-backs.