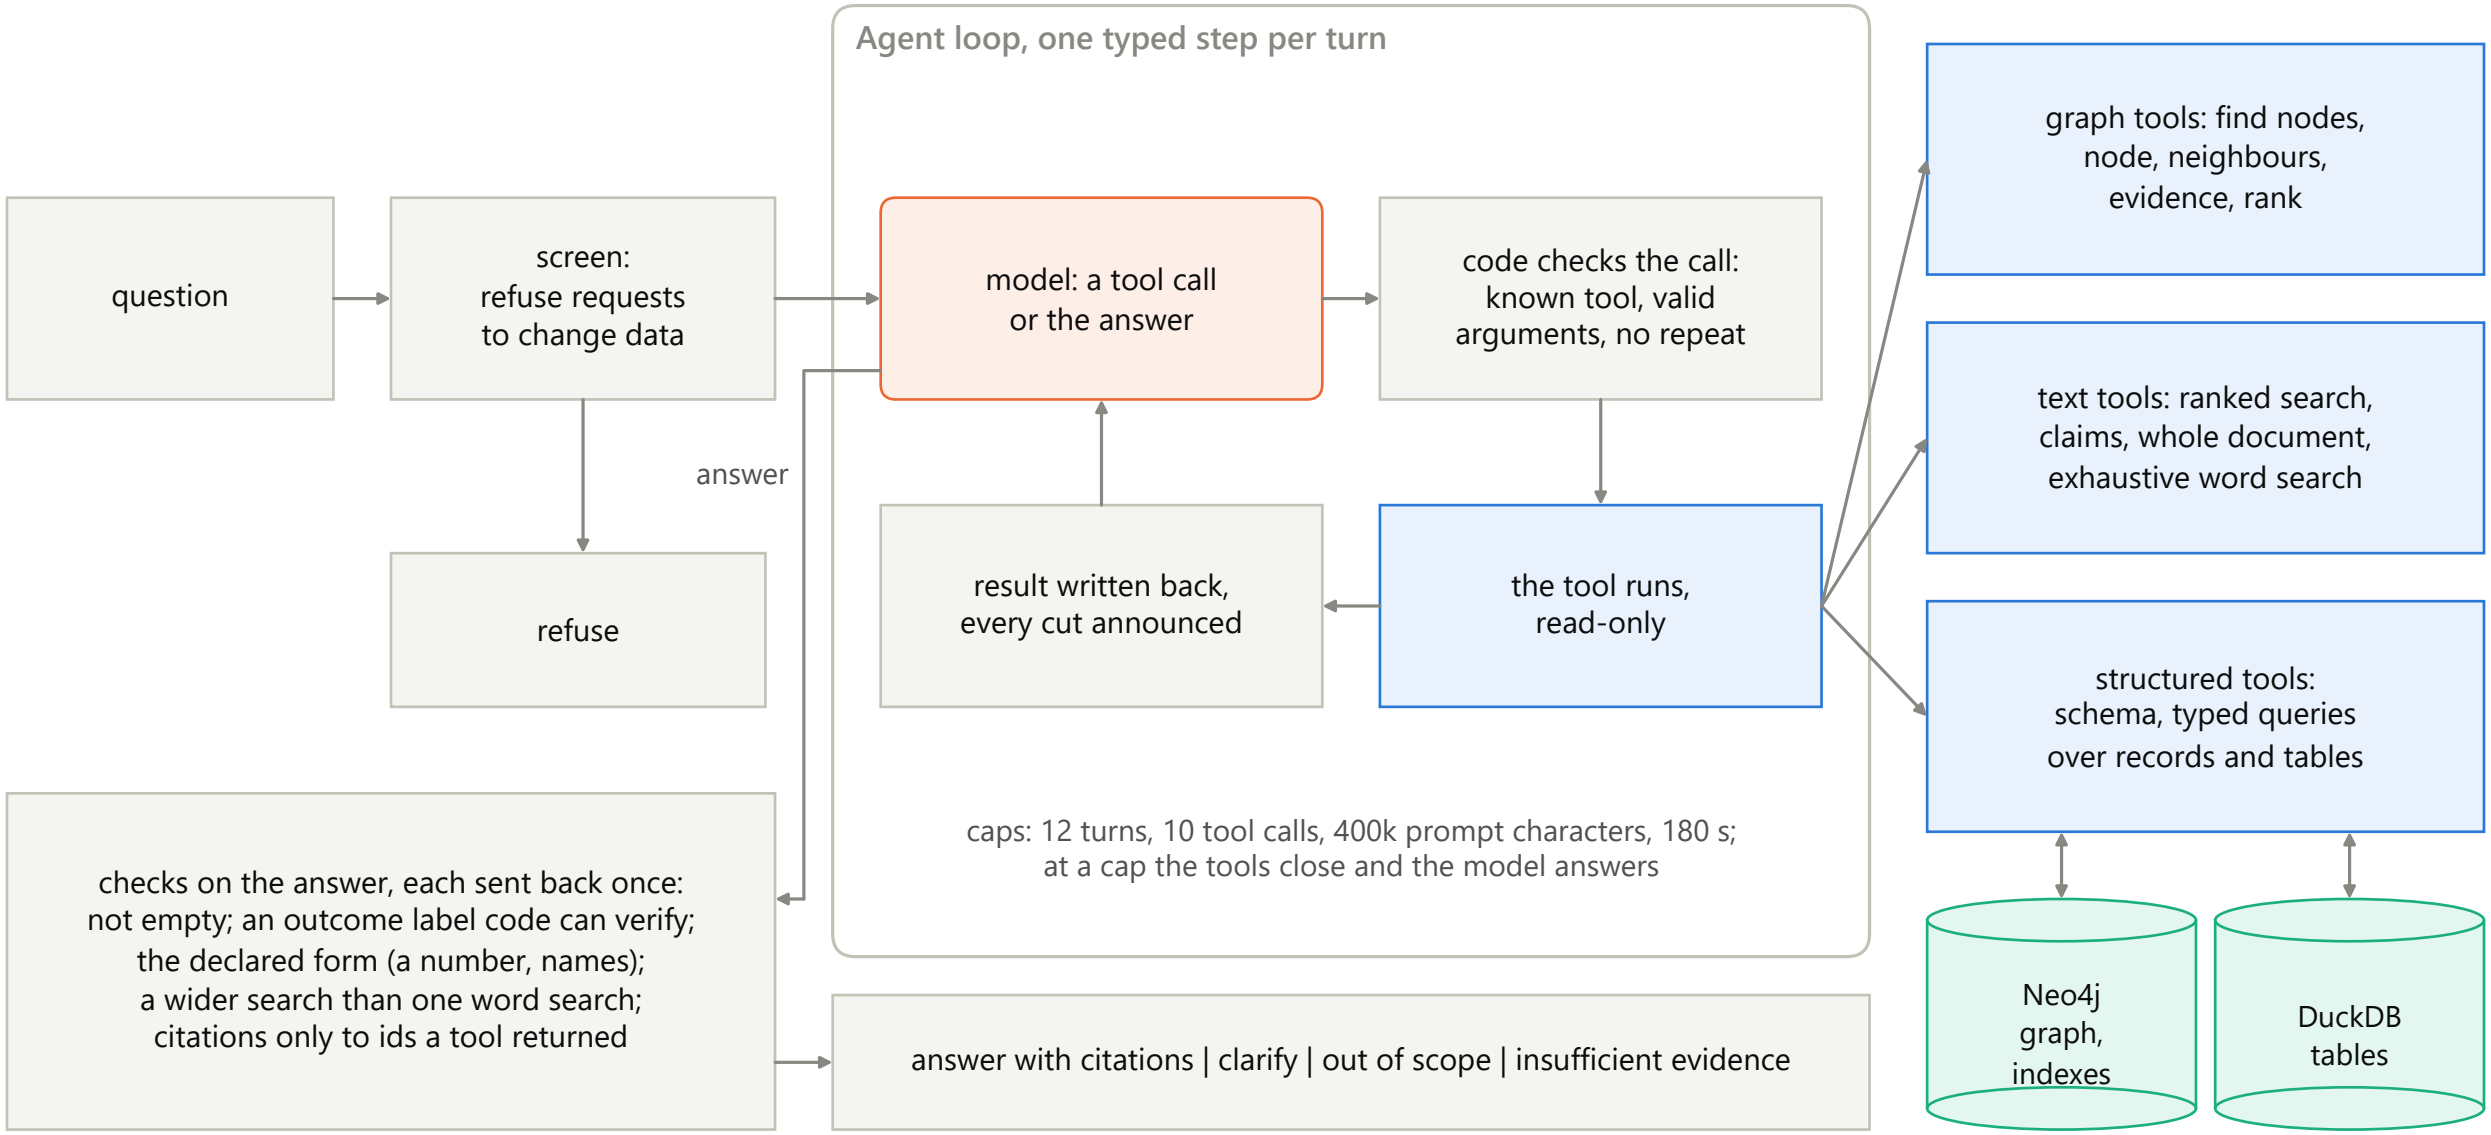
Agent loop, one typed step per turn
question
screen:
refuse requests
to change data
refuse
answer
model: a tool call
or the answer
code checks the call:
known tool, valid
arguments, no repeat
result written back,
every cut announced
the tool runs,
read-only
graph tools: find nodes,
node, neighbours,
evidence, rank
text tools: ranked search,
claims, whole document,
exhaustive word search
structured tools:
schema, typed queries
over records and tables
Neo4j
graph,
indexes
DuckDB
tables
caps: 12 turns, 10 tool calls, 400k prompt characters, 180 s;
at a cap the tools close and the model answers
checks on the answer, each sent back once:
not empty; an outcome label code can verify;
the declared form (a number, names);
a wider search than one word search;
citations only to ids a tool returned
answer with citations | clarify | out of scope | insufficient evidence