7 of 200 benchmark answers are wrong by either score; every one wrong by meaning involves the agent's reasoning
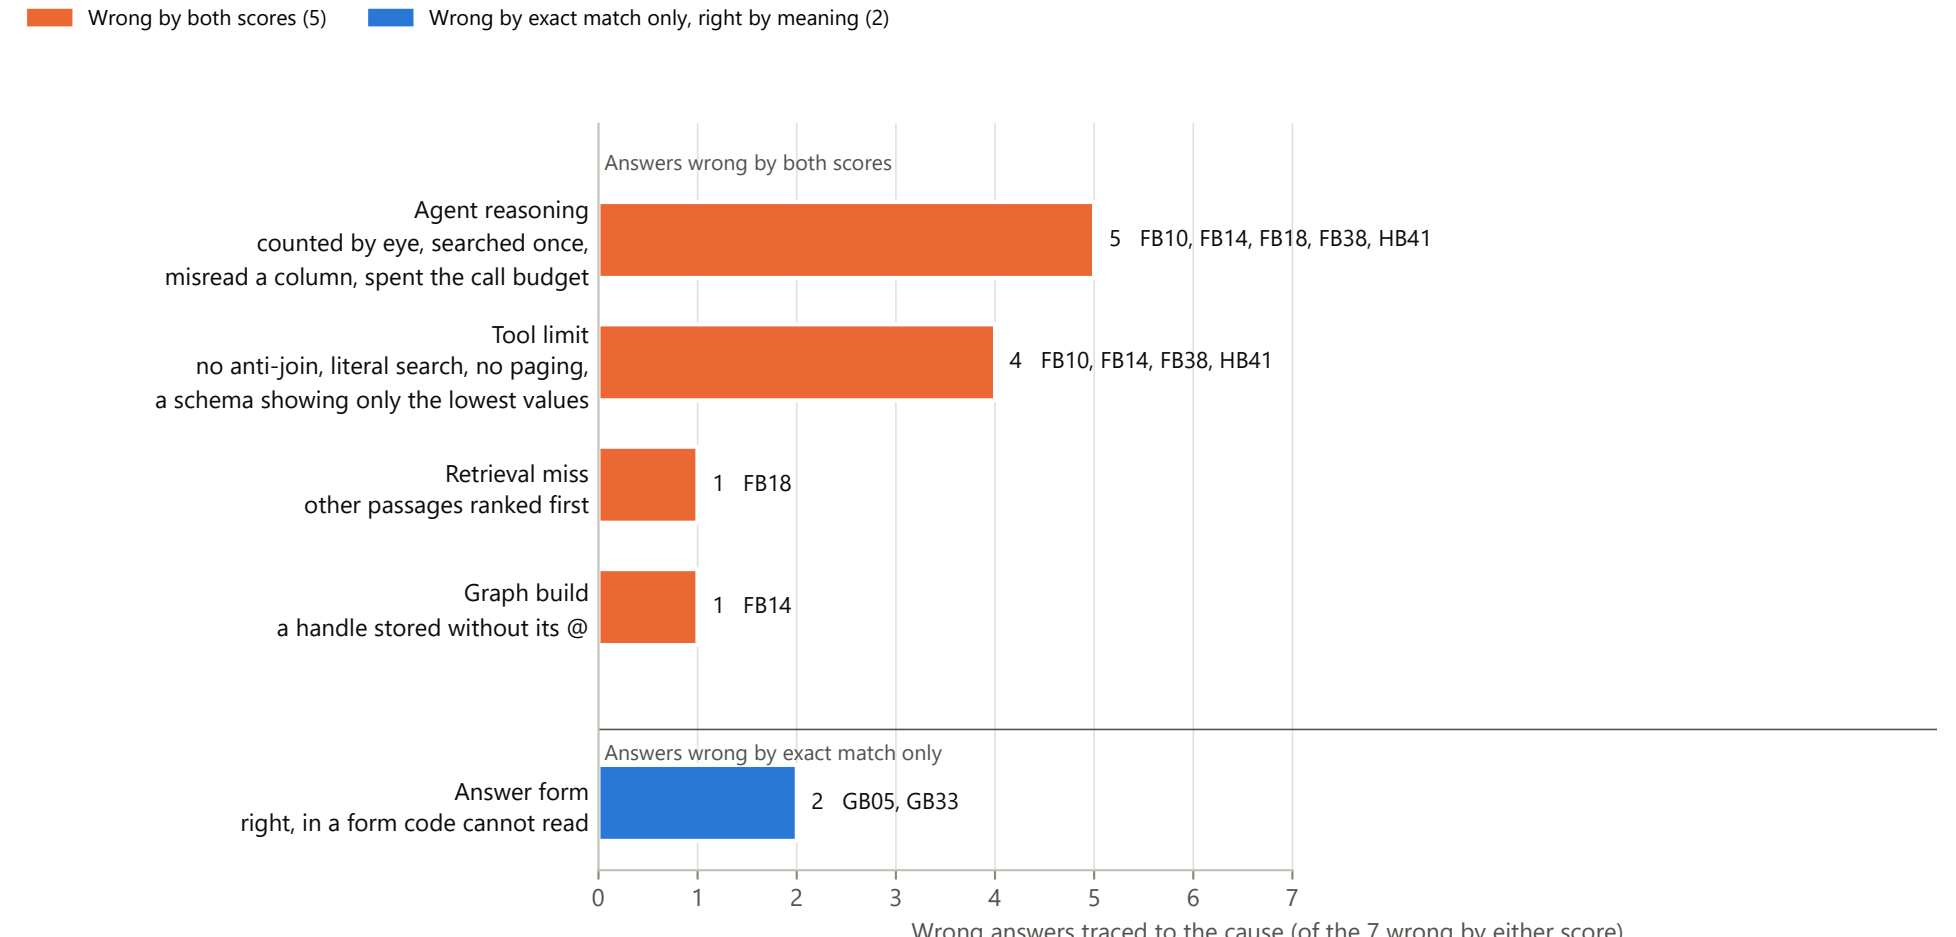
Four datasets: NTSB, furniture, held-out, generality. R157: the wrong answers, each traced to its causes; a question can have several, so the bars add up to more than 7.
Question ids at each bar: FB furniture, HB held-out, GB generality; no NTSB answer was wrong.